



Deadlocks

What is a *deadlock*?

- **Deadlock** -- *permanent* blocking of a set of processes that compete for system resources
- **No efficient solution** to the deadlock problem in the general case.
- In deadlock
 - ◆ a set of **processes are in a wait** state, because each process is waiting for a resource that is held by some other waiting process
- **All deadlocks involve conflicting resource needs by two or more processes**

Classification of resources—I

- Resources grouped into two classes:
 - ◆ **Reusable**: something that can be safely used by one process at a time and is ***not depleted*** by that use. Processes obtain resources that they later release for reuse by others.
E.g., CPU, memory, specific I/O devices, or files.
 - ◆ **Consumable**: these can be ***created and destroyed***. When a resource is acquired by a process, the resource ceases to exist.
E.g., interrupts, signals, or messages.

Classification of resources—II

- Another classification of resources:
 - ◆ **Preemptable**: these can be taken away from the process owning it with no ill effects (needs save/restore).
E.g., memory or CPU.
 - ◆ **Non-preemptable**: cannot be taken away from its current owner without causing the computation to fail.
E.g., printer or floppy disk.
- Deadlocks mostly occur when sharing **reusable** and **non-preemptable** resources.

Conditions for deadlock

- Four conditions that must hold for a deadlock to be possible:
 - ◆ **Mutual exclusion:** processes require exclusive control of its resources (not sharing)
 - ◆ **Hold and wait:** process may wait for a resource while holding others
 - ◆ **No preemption:** process will not give up a resource until it is finished with it
 - ◆ **Circular wait:** each process in the chain holds a resource requested by another

Conditions for deadlock

- ***Also, a process cannot be reset to an earlier state where resources not held***
- Conditions 1—3 are necessary, but not sufficient. All 4 are needed

Solving deadlocks

- If a necessary condition is prevented a deadlock cannot occur. For example:
 - ◆ Systems with only shared resources cannot deadlock.
 - Negates *mutual exclusion*.
 - ◆ Systems that abort processes which request a resource that is in use.
 - Negates *hold and wait*.

Solving deadlocks (contd.)

- If a necessary condition is prevented a deadlock cannot occur. For example:
 - ◆ Preemptions may be possible if a process does not use its resources until it has acquired all it needs.
 - Negates *no preemption*.
 - ◆ Systems that detect or avoid deadlocks.
 - Prevents *cycle*.
 - ◆ Transaction processing systems provide checkpoints so that processes may back out of a transaction.
 - Negates *irreversible process*.

Resource allocation graphs

Set of Processes


$P = \{P_1, P_2, \dots, P_n\}$

Set of Resources


$R = \{R_1, R_2, \dots, R_m\}$

Some resources come in multiple units.

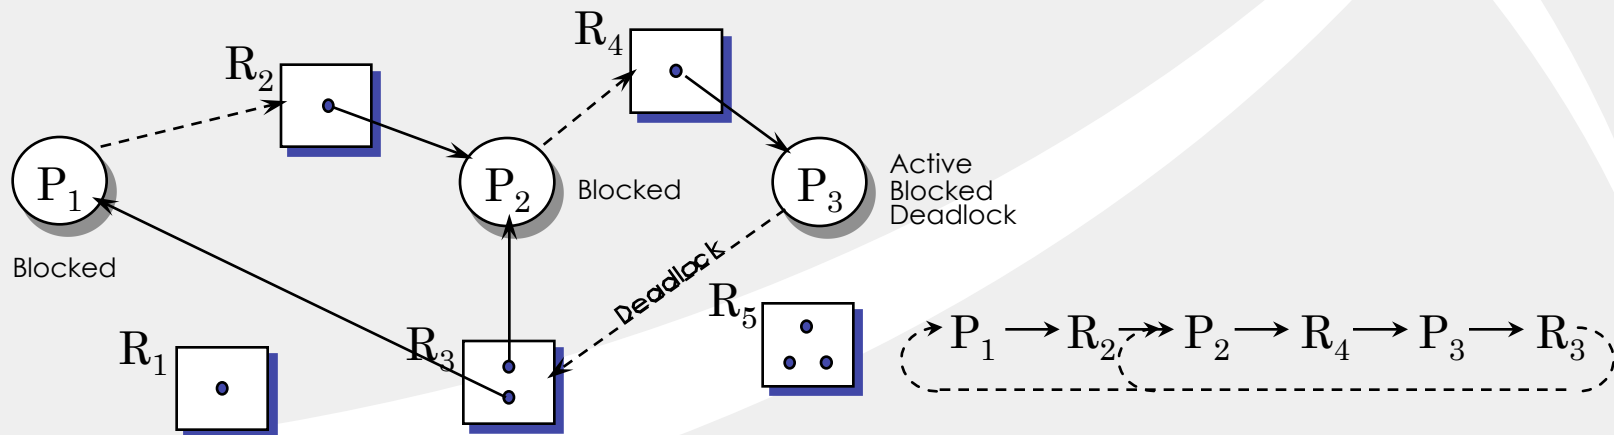
R_j has 2 units




Process P_i waits for (has requested) R_j



Resource R_j has been allocated to P_i



Cycle is necessary, but ...



Multiple resource unit case:


No Deadlock—yet!

Because, either P₂ or P₄ could relinquish a resource allowing P₁ or P₃ (which are currently blocked) to continue. P₂ is still executing, even if P₄ requests R₁.


... a knot is required

- Cycle is a *necessary condition* for a deadlock. With multiple unit resources—*not sufficient*.
- A *knot* must exist—a cycle with no non-cycle outgoing path from any involved node.
- At the moment assume that:
 - ◆ a process *halts* as soon as it waits for one resource, and
 - ◆ processes can wait for only *one* resource at a time

Further requests



(a) If P_4 requests R_1 no deadlock since P_2 may release R_1 , allowing P_4 to complete.



(b) If P_2 requests R_2 : Deadlock—Cycle—Knot.

No active processes to release resources.

Strategies for deadlocks

- In general, four strategies are used for dealing with deadlocks:
 - ◆ **Ignorance:** pretend there is no problem at all.
 - ◆ **Prevention:** design a system in such a way deadlock is excluded a priori.
 - ◆ **Avoidance:** make a decision dynamically checking whether a request will, if granted, potentially lead to a deadlock or not.
 - ◆ **Detection:** let the deadlock occur and detect when it happens, and take some action to recover after the fact.

Dealing with Deadlocks

- Different people react to this strategy in different ways:
 - ◆ **Mathematicians:** find deadlock totally unacceptable, and say that it must be prevented at all costs.
 - ◆ **Engineers:** ask how serious it is, and do not want to pay a penalty in performance and convenience.
- UNIX approach
 - ◆ ignore the problem
 - ◆ Prevention price is high – inconvenient restrictions

Deadlock prevention

- Deadlock prevention is to design a system to exclude the possibility of a deadlock
 - ◆ *indirect methods* -- prevent the occurrence of one of the necessary conditions listed earlier.
 - ◆ *direct methods* -- prevent the occurrence of a circular wait condition.
- Deadlock prevention strategies are very **conservative** -- solve the problem of deadlock by limiting access to resources and by imposing restrictions on processes

More on deadlock prevention

◆ Mutual exclusion

- In general, this condition cannot be disallowed.

◆ Hold-and-wait

- The hold and-wait condition can be prevented by requiring that a process request all its required resources at one time, and blocking the process until all requests can be granted simultaneously.


More on deadlock prevention(cntd)

◆ No preemption

- One solution is that if a process holding certain resources is denied a further request, that process must release its unused resources and request them again, together with the additional resource.

◆ Circular Wait

- The circular wait condition can be prevented by defining a linear ordering of resource types. If a process has been allocated resources of type R , then it may subsequently request only those resources of types following R in the ordering.



Deadlock avoidance

- Allows the necessary conditions
- Makes judicious resource allocation choices to ensure a deadlock-free system
 - ◆ System evaluates current request and denies it if granting it will lead to potential deadlock
 - ◆ Requires knowledge of future requests
- Deadlock avoidance approaches:
 - ◆ Resource trajectories
 - ◆ Safe/unsafe states
 - ◆ Dijkstra's Banker's algorithm

Deadlock Avoidance

Requires that the system has some additional *a priori* information available

- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes


Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the systems such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
- That is:
 - ◆ If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
 - ◆ When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - ◆ When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Basic Facts

- If a system is in safe state \Rightarrow no deadlocks
- If a system is in unsafe state \Rightarrow possibility of deadlock
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

Safe, Unsafe, Deadlock State




Avoidance Algorithms

- Single instance of a resource type
 - ◆ Use a resource-allocation graph
- Multiple instances of a resource type
 - ◆ Use the banker's algorithm


Resource-Allocation Graph Scheme

- **Claim edge** $P_i \rightarrow R_j$ indicated that process P_i may request/wait resource R_j ; represented by a dashed line
- Claim edge converts to request/wait edge when a process requests a resource
- Request edge converted to an assignment/allocation edge when the resource is allocated to the process
- When a resource is released by a process, assignment/allocation edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system

Resource-Allocation Graph



Unsafe State In Resource-Allocation Graph



Resource-Allocation Graph Algorithm

- Suppose that process P_i requests a resource R_j
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

Banker's Algorithm

- Multiple instances
- Each process must a priori claim maximum use
- When a process requests a resource it may have to wait
- When a process gets all its resources it must return them in a finite amount of time

Banker's algorithm—*definitions*

Assume N Processes $\{P_i\}$
 M Resources $\{R_j\}$

Availability vector **Avail_j**, units of each resource (initialized to maximum, changes dynamically).

Let [**Max_{ij}**] be an $N \times M$ matrix.

Max_{ij} = L means Process P_i will request at most L units of R_j .

[**Hold_{ij}**] Units of R_j currently held by P_i

[**Need_{ij}**] Remaining need by P_i for units of R_j

Need_{ij} = **Max_{ij}** - **Hold_{ij}**, for all i & j

Banker's algorithm—*resource request*

- At any instance, P_i posts its request for resources in vector REQ_j .
- **Step 1:** verify that a process matches its needs.
■ *if* $REQ_j > Need_{ij}$ *abort—error, impossible*
- **Step 2:** check if the requested amount is available.
■ *if* $REQ_j > Avail_j$ *goto Step 1— P_i must wait*
- **Step 3:** provisional allocation.
■ $Avail_j = Avail_j - REQ_j$
■ $Hold_{ij} = Hold_{ij} + REQ_j$
■ $Need_{ij} = Need_{ij} - REQ_j$
■ *if* $isSafe()$ *then* grant resources—system is
safe
- *else* cancel allocation; *goto Step 1— P_i must wait*

Banker's algorithm—*isSafe*

Find out whether the system is in a safe state.

Work and **Finish** are two temporary vectors. *Temporary availability vector that will be modified in this function*

Step 1: initialize.

$\text{Work}_j \leftarrow \text{Avail}_j$ for all j ; $\text{Finish}_i = \text{false}$ for all i .

Step 2: find a process P_i such that

$\text{Finish}_i = \text{false}$ and $\text{Need}_{ij} \leq \text{Work}_j$
if no such process, *goto Step 4.*

Step 3: $\text{Work}_j = \text{Work}_j + \text{Hold}_{ij}$

$\text{Finish}_i = \text{true}$

goto Step 2.

Process i needs are < than available resources j
Process i is done, See if other processes need checking

Step 4: if $\text{Finish}_i = \text{true}$ for all i

then return true—yes, the system is safe

else return false—no, the system is NOT safe

P_i, REQ_j

Banker's algorithm—*what is safe?*

■ Safe with respect to some resource allocation.

◆ very safe

$NEED_i \leq AVAIL$ for all Processes P_i ← *resource needs are available for all processes*
Processes can run to completion in any order.

◆ safe (but take care)

$NEED_i > AVAIL$ for some P_i

$NEED_i \leq AVAIL$ for at least one P_i such that

There is at least one correct order in which the processes may complete their use of resources.

◆ unsafe (deadlock inevitable)

$NEED_i > AVAIL$ for some P_i

$NEED_i \leq AVAIL$ for at least one P_i

But some processes cannot complete successfully.

◆ deadlock


$NEED_i > AVAIL$ for all P_i ← *No process can obtain the required resources to run*

Processes are already blocked or will become so as they request a resource.

Example—safe allocation *cont.*

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	2 ₃	3 ₂	F	2	2 ₁
P ₂	4	1	3	F		
P ₃	2	1	1	F		

Assume P₁ acquires one unit.



Example—safe allocation

	Max	Hold	Need	Finish	Avail	Work
P_1	5	2 ₃	3 ₂	F	2	2 ₁
P_2	4	1	3	F		
P_3	2	1	1	F		

For simplicity, assume that all the resources are identical.

Assume P_1 acquires one unit. *Very safe?* **No!** $\text{Need}_2 > 2 \leftarrow \text{avail}$

Safe? Let us see with the safe/unsafe algorithm...

$i = 1$; does P_1 agree with **Step 2**? No. $(\text{Finish}_i = \text{false and } \text{Need}_{ij} \leq \text{Work}_j) ?$


$i = 2$; does P_2 agree with **Step 2**? No.

$i = 3$; does P_3 agree with **Step 2**? Yes.

Example—safe allocation *cont.*

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	3	2	F	1	1
P ₂	4	1	3	F		0
P ₃	2	2	0	F		

P₃ can acquire the last unit and finish.




Work = Work + Hold₃; Finish₁ = T
Work = 0 + 2 = 2

Example—safe allocation

cont.

	Max	Hold	Need	Finish	Avail	Work
P_1	5	5	0	F	2	2 ₀
P_2	4	1	3	F		
P_3	2	0	0	T		

Then, P_1 can acquire two more units and finish.




($Finish_i = \text{false}$ and $Need_{ij} \leq Work_j$)?

$i = 1$; does P_1 agree with **Step 2**? Yes. $Work = Work + Hold_1$; $Finish_1 = T$
 $Work = 0 + 5 = 5$

Example—safe allocation *cont.*

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	0	0	T	5	2
P ₂	4	1 4	3 0	F		
P ₃	2	0	0	T		

Finally, P₂ can acquire three more units and finish.



(Finish_i = false and Need_{i,j} ≤ Work_j)?

i = 2; does P₂ agree with **Step 2**? Yes. Work = Work+Hold₂; Finish₂ = T

Example—unsafe allocation

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	2	3	F	2	2 ₁
P ₂	5	1 ₂	4 ₃	F		
P ₃	2	1	1	F		

Assume P₂ acquires one unit.

As before, P₃ can finish and release its resources.

BUT...

i = 1; does P₁ agree with **Step 2**? No.

i = 2; does P₂ agree with **Step 2**? No.

i = 3; does P₃ agree with **Step 2**? Yes. **Work** = **Work**+**Hold**₃=2; **Finish**₂ = **T**

Any more unfinished P_i? Yes.


P₁ and P₂ cannot finish. Therefore *unsafe*.

Example—unsafe allocation

cont.

	Max	Hold	Need	Finish	Avail	Work
P ₁	5	2	3	F	2	2
P ₂	5	2	3	F		
P ₃	2	0	0	T		

NOW...



Deadlock detection

- This technique
 - ◆ does not attempt to prevent deadlocks;
 - ◆ instead, it lets them occur.
- The system
 - ◆ detects when this happens, by periodically running an algorithm to detect a **circular wait** condition,
 - ◆ then takes some action to recover after the fact.
- With deadlock detection, requested resources are granted to processes whenever possible.

Deadlock detection

cont.

- A check for deadlock can be made
 - ◆ as frequently as for each resource request, or
 - ◆ less frequently, depending on how likely it is for a deadlock to occur.
- Checking at each resource request has two advantages:
 - ◆ it leads to early detection, and
 - ◆ the algorithm is relatively simple because it is based on incremental changes to the state of the system.
- On the other hand,
 - ◆ such frequent checks consume considerable processor time.

Recovering from deadlocks

- Once the deadlock algorithm has successfully detected a deadlock, some strategy is needed for recovery. There are various ways:

- ◆ Recovery through *Preemption*

In some cases, it may be possible to temporarily take a resource away from its current owner and give it to another.

- ◆ Recovery through *Rollback*

If it is known that deadlocks are likely, one can arrange to have processes *checkpointed* periodically. For example, can undo transactions, thus free locks on database records.

- ◆ Recovery through *Termination*

The most trivial way to break a deadlock is to kill one or more processes. One possibility is to kill a process in the cycle. Warning! *Irrecoverable losses may occur, even if this is the least advanced process.*

Summary of strategies

<i>Principle</i>	<i>Resource Allocation Strategy</i>	<i>Different Schemes</i>	<i>Major Advantages</i>	<i>Major Disadvantages</i>
<i>DETECTION</i>	<ul style="list-style-type: none"> • Very liberal; grant resources as requested. 	<ul style="list-style-type: none"> • Invoke periodically to test for deadlock. 	<ul style="list-style-type: none"> • Never delays process initiation. • Facilitates on-line handling. 	<ul style="list-style-type: none"> • Inherent preemption losses.
<i>PREVENTION</i>	<ul style="list-style-type: none"> • Conservative; under-commits resources. 	<ul style="list-style-type: none"> • Requesting all resources at once. • Preemption • Resource ordering 	<ul style="list-style-type: none"> • Works well for processes with single burst of activity. • No preemption is needed. • Convenient when applied to resources whose state can be saved and restored easily. • Feasible to enforce via compile-time checks. • Needs no run-time computation. 	<ul style="list-style-type: none"> • Inefficient. • Delays process initiation. • Preempts more often than necessary. • Subject to cyclic restart. • Preempts without immediate use. • Disallows incremental resource requests.
<i>AVOIDANCE</i>	<ul style="list-style-type: none"> • Selects midway between that of detection and prevention. 	<ul style="list-style-type: none"> • Manipulate to find at least one safe path. 	<ul style="list-style-type: none"> • No preemption necessary. 	<ul style="list-style-type: none"> • Future resource requirements must be known. • Processes can be blocked for long periods.